

# GENERAL PROBLEM SOLVING WITH SEARCH ALGORITHMS

# 8-PUZZLE PROBLEM SOLVING

*Input:*

1	2	3
	4	5
7	8	6

*Goal:*

1	2	3
4	5	6
7	8	

# 8-PUZZLE PROBLEM SOLVING

*Input:*

1	2	
4	5	3
7	8	6

*2 possible moves:*

1		2
4	5	3
7	8	6

1	2	3
4	5	
7	8	6

1	2	3
4	5	6
7	8	

*Goal*

# 8-PUZZLE PROBLEM SOLVING

*Input:*

1	2	
4	5	3
7	8	6

*Branching factor,  $b=2$*

1		2
4	5	3
7	8	6

1	2	3
4	5	
7	8	6

$b=3$

$1, R$

$2, L$

$5, U$

$b=3$

$3, D$

$5, R$

$6, U$

*Goal*

1	2	3
4	5	6
7	8	

# Problem Solving = Graph Search

- Search algorithm creates a *Search Tree*
- *Branching Factor* may increase or decrease, *for the above*: 2, 3, 4
- **Tree Search:** Same pattern may be repeated!
  - That may mean looping
- **Really, Graph Search:** use memory to remember **visited** nodes

# Problem Solving = Graph Search

- Note, in AI it is being <<dynamically>> performed: *Generate-and-Test*
  - *Test* in each iteration, if solution is reached
  - Search tree is dynamically getting *generated* during search
  - *Generate-test (AI-search) = Map-reduce (LISP) = Map-reduce (Hadoop-BigData)*
- *AI search*: The Graph may NOT be *statically* available as input
- Some input board may not even have any solution!

# Problem Solving = Graph Search

- **Problem Solving in AI is often (always!): Graph Search**
  - It is mostly graph search, but often we treat it as a tree search,
  - ignoring, repeat visit of nodes
    - because memorizing & checking for past-visited nodes is too expensive!
- Each node is a “state”: *State space of nodes is searched*
- *Control types:*
  - *Breadth First Search (BFS)*
  - *Depth First Search (DFS)*

# Sliding Puzzle Solving with BFS/DFS

*Input or Problem instance*

1	2	
4	5	3
7	8	6

*Goal*

1	2	3
4	5	6
7	8	


- *Think of input data structure.*
- **Input Size  $n$ =**
  - *8-Puzzle: 3x3 slots,  $n=9$*
  - *15-Puzzle: 4x4 slots,  $n=16$*
  - ...
- *Problem Complexity= Total number of feasible steps over a search tree*
  - *What are the steps: **Left/Right/Up/Down***
  - *Complexity as  $O(f(n))$*
  - *#Steps:  **$O(b^n)$** , for average branching factor  **$b$** :*
    - *Branching means, how many options available at each step*



# AI Problem Solving

- *NP-hard problems: likely to be  $O(k^n)$ , for some  $k > 1$ , exponential*
  - *Clever, efficient algorithms exist, but worst case remains exponential*
  - *AI is often about efficiently solving with clever algorithm!*

# BFS: 8-PUZZLE PROBLEM SOLVING



*What is the typical data structure for BFS? Queue or Stack?*

# Breadth First Search

*Algorithm BFS(s): s start node*

*0. Initialize a Q with s;*

*1. v = Pop(Q); // BFS is Queue based algorithm*

*2. If v is “goal” return success;*

*3. mark node v as visited; // done!*

*// absent in “tree” search mode*

*4. operate on v; // e.g., evaluate if it is the goal node, if so, return*

*5. for each node w accessible from node v do*

*6.       if w is not marked as visited then    // with generate-and-test this may not be time consuming*

*7.               Push w at the back of Q;*

*end for;*

*End algorithm.*

*Code it!*

# DFS: 8-PUZZLE PROBLEM SOLVING

*Input:*

1	2	
4	5	3
7	8	6

*Red and green arrows are control flow*

1		2
4	5	3
7	8	6

1	2	3
4	5	
7	8	6

$b=3$

$1, R$

$2, L$

$5, U$

$b=3$

$3, D$

$5, R$

$6, U$

*Until leaf node,  
no more children*

*Goal*

# Depth First Search

**Algorithm** *DFS(v)* // recursive

*1.mark node v as visited; // done!*

*// absent in “tree” search mode*

*2.If v is “goal” return success;*

*3.operate on v ; // e.g., evaluate if it is the goal node, if so, return*

*4.for each node w accessible from node v do*

*5. if w is not marked as visited then // again, absent in “tree” search mode*

*6. **DFS(w)**; // an iterative version of this  
//will maintain its own stack*

*end for;*

**End algorithm.**

## Driver algorithm

*Input:* Typical Graph search input is a Graph  $G$ : (nodes  $V$ , arcs  $E$ )

*// Here we have: a board as a node  $p$ , and operators for modifying board correctly*

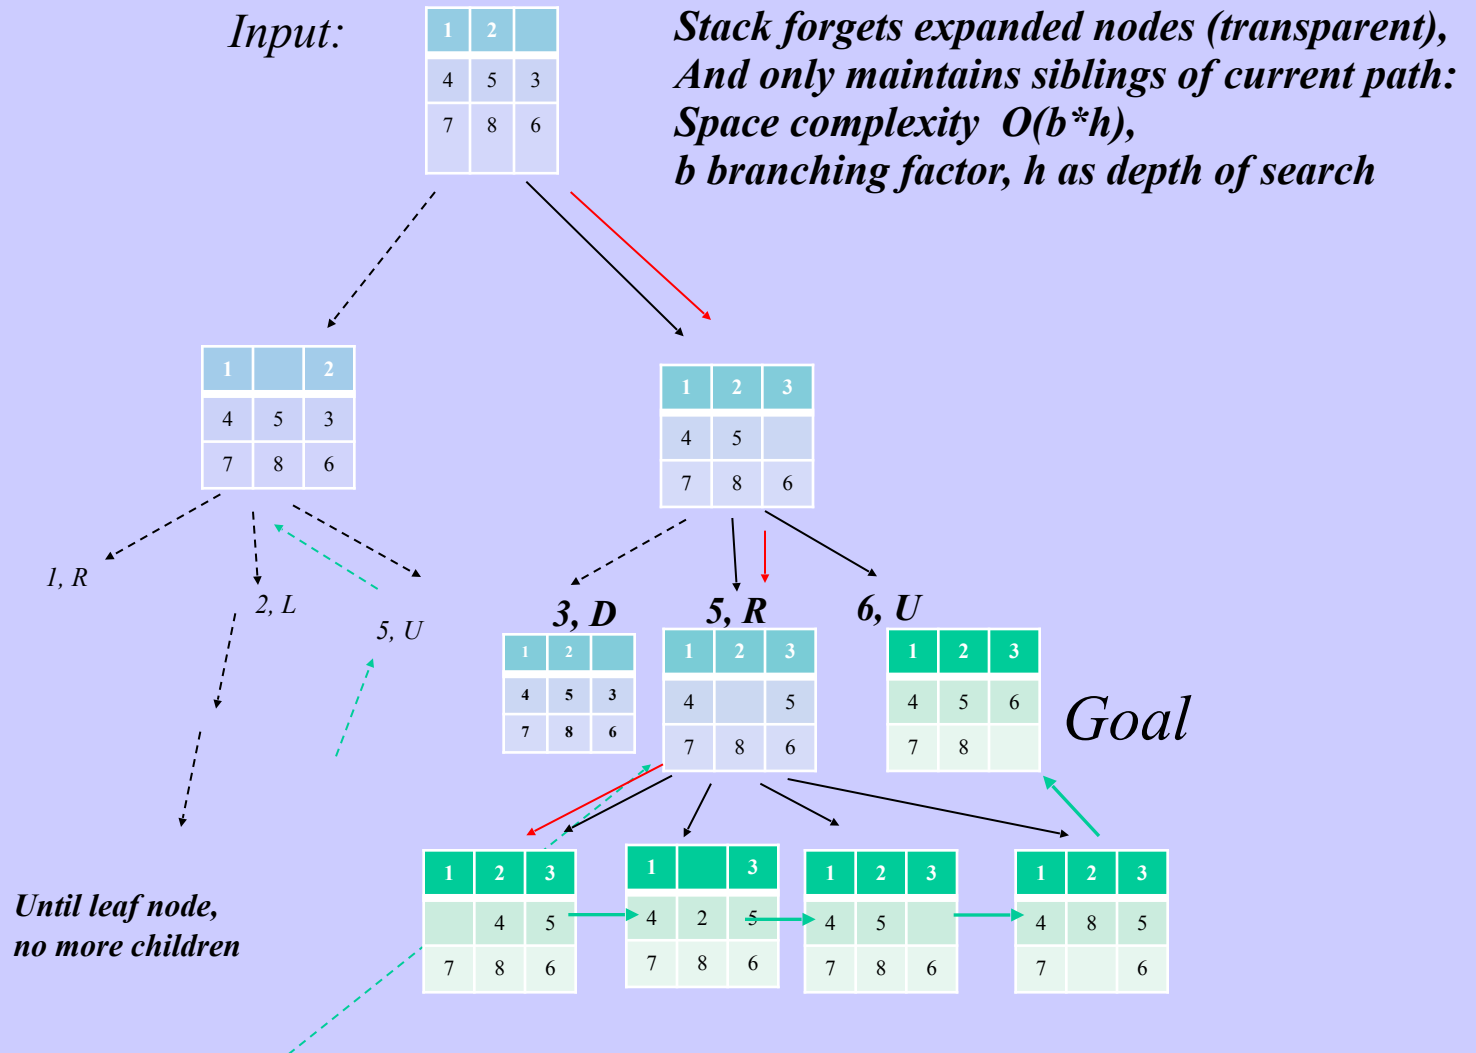
*// i.e, we generate the graph as we go: “Generate and Test”*

*Output:* Path to a goal node  $g$ , computer may display the next node on the path

**call *DFS(p)*; // input board  $p$  as a node in search tree  $G$**

**End.**

# DFS: 8-PUZZLE PROBLEM SOLVING



# BFS vs DFS

- **BFS:**


- **Memory intensive: ALL children are to be stored from all levels** ☹
  - *At least all nodes of the last level are to be stored, and that grows fast too!*
  - *Also, you cannot do the above if the path from start node to goal is needed!*
- **If goal exists, BFS is guaranteed to find it: complete algorithm**
  - *(systematically finds it, level by level)* ☺
- **If goal is nearby (at a shallow level), BFS quickly finds it** ☺
- **Memory intensive, all nodes in memory** ☹
- **Queue for implementation**

- **DFS:**


- **Infinite search may happen, in the worst case** ☹
  - *May get stuck on an **infinite depth** in a branch,*
  - *Even if the goal may be at a shallow level on a different branch*
- **Linear memory growth, depth-wise** ☺ **WHY?**
  - *but, go to the point number 1 above → memory may explode for large depth!*
- **Stack for implementation (equivalent to recursive implementation)**

# BFS vs DFS

- *Time complexity: Worst case for both, all nodes searched:  $O(b^d)$   
    *b* branching factor, *d* depth of the goal*
- *Memory:*
  - *BFS  $O(b^d)$  remember all nodes*




- **DFS**  $O(bd)$ , only one set of children at each level, up to goal depth
  - But the depth may be very large, up to infinity ☹️
  - DFS **<forgets>** previously explored branches 😊





Depth Limited Search (DLS):  
*To avoid infinite search of DFS*

- Stop DFS at a *fixed* depth  $l$ , no matter what
- Goal, may NOT be found: *Incomplete Algorithm*
  - If goal depth  $d > l$
- Why DLS? To avoid getting stuck at infinite (read: large) depth
- Time:  $O(b^l)$ , Memory:  $O(bl)$



# Iterative Deepening Search (IDS)


- ***Repeatedly*** stop DFS at a fixed depth  $l$  (i.e., DLS), **AND**
  - If goal is not found,
  - then **RESTART** from start node again, with  $l = l+1$
- ***Complete Algorithm***
  - Goal <<will>> be found when  $l = d$ , goal depth
- Isn't repetition very expensive?
- **Time:**  $O(b^1 + b^2 + b^3 + \dots + b^d) = O(b^{d+1})$ , as opposed to  $O(b^d)$ ,
  - For  $b=10$ ,  $d=5$ , this is 111,000 to 123,450 increase, 11%
- **Memory:**
  - **IDS vs DFS:** same  $O(bd)$
  - **IDS vs BFS:**  $O(bd)$  vs  $O(b^d)$

## **Informed Search (Heuristics Guided Search)**

# Informed Search (Heuristics Guided Search, or “AI search”)

- Available is a “heuristic” function  $f(n)$  node  $n$ , to chose a **best** Child node from alternative nodes
- Example: **Guessed** distance from goal
  - *You can see the Eiffel tower (goal), move towards it!*
- Best\_child = Child  $w$  with **minimum**  $f(w)$ 
  - We will use  $h(w)$  for node  $w$  as guessed distance to goal
- AI-searches are **optimization** problems
- Both BFS and DFS may be **guided**: Informed search

# Best First Search



## *Algorithm ucfs(v)*

1. *Enqueue start node  $s$  on a min-priority queue  $Q$ ; // distance  $g(s) = 0$*
2. *While  $Q \neq \text{empty}$  do*
3.      $v = \text{pop}(Q)$ ;
4.     *If  $v$  is a goal, return success;*
5.     *Operate on  $v$  (e.g., display board);*
6.     *For each child  $w$  of  $v$ ,*
7.     *Insert  $w$  in  $Q$  such that  $\text{cost}(w)$  is the lowest;*  
      *//  $\text{cost}(w)$  is path cost from the child node  $w$  to a goal node,*  
      *// replacing  $\text{cost}(w)$  with  $g(w)$ , path-cost of  $w$  from source, makes it*

## *Dijkstra's algorithm*

*//  $Q$  is a priority-queue to keep lowest cost-node in front*

# Uniform Cost Search:


*Breadth First with weights rather than levels*

*Algorithm ucs(v)*

1. Enqueue start node  $s$  on a min-priority queue  $Q$ ; // distance  $h(s) = 0$
2. While  $Q \neq \text{empty}$  do
3.      $v = \text{pop}(Q)$ ; // **lowest path-cost node**
4.     If  $v$  is a goal return success;
5.     operate on  $v$  (e.g., display board);
6.     For each child  $w$  of  $v$ , if  $w$  is not-visited before,
7.         Insert  $w$  in  $Q$  such that  $g(w)$  **is the lowest**; // use priority-que or heap for  $Q$

*Status of  $Q$  on BFS:*  
(when  $g$  is not used)

$Q: 30$   
 $Q: 5, 10$   
 $Q: 10, 11, 7$   
 $Q: 11, 7, 12, G$   
 $Q: 7, 12, G, x$   
 $Q: 12, G, x, \dots$   
 $Q: G, x, \dots$   
*Success*



*Status of  $Q$  on ucbs:*

$Q: 30$   
 $Q: 5, 10$   
 $Q: 7, 10, 11$  // because it is a heap  
 $Q: 10, 11$   
 $Q: 11, 12, \dots$       $10 \rightarrow G$ : **Success**



*Consider mouse/robot in a maze*

# Best First Search:

*Find shortest path to a goal*

*Similar to Dijkstra's shortest path-finding algorithm (from source)*

**In a graph, the cost of a node may be updated via a different path**



*Status of Q:*

*Q: 80, 99*

*Q: 99, 177*

*310, Success, but continue until finished*

*Q: 177, 278*






*Q: 278*

*278, Success, and finished all nodes*

*We may stop if just finding a Goal is enough,  
And shortest path is not needed,  
e.g., in typical 8-puzzle problem*

# Animation of Best First Search

*transparent nodes below are no longer in the Priority Queue*





## Two properties of search algorithms

- **Completeness:** does it find a **goal** if it **exists** within a **finite** size path?
- **Optimality:** does it find the **shortest** path-distance to a goal?
  - **Note:** not to confuse with *optimal-algorithm*, as in complexity theory:  
lowest time in Omega  $\Omega$  notation
  - Our “optimal” is *optimal-search* algorithm, in AI
  - “Completeness” is somewhat similar in both the context

# Greedy search: A\* Search

- “Heuristic” function  $f(n) = g(n) + h(n)$   
 $g(n)$  = current distance from start to current node  $n$   
 $h(n)$  = a *guessed distance* to goal
- *Best-first strategy:*
  - *Out of all frontier nodes pick up  $n'$  that has minimum  $h(n')$*
  - *Use heap*

*Mouse/robot can see a pole at the exit gate*

# Greedy Depth-first search: $A^*$ Search


- **Example:**
  - $g(n)$ : *how far have you traveled*
  - $h(n)$ : *how far does the Eiffel tower appear to be!*
  - *Go toward smallest  $f = g+h$* 
    - *least  $f =$  least path to goal*
    - *Why least path, rather than any path?*
    - *Because: a larger path may be infinite!*

## A\* Search is *Best-first with $f$*

- $f$  is “*correct*”, iff DFS is guided toward the goal correctly
  - “**Admissible**” heuristic:  
 $f$  is  $\leq$  **actual  $g(goal)$** , in **minimize** problem
- Say,  $f(n) = 447$ , then  
there should be **no path** to goal via  $n$  costing 440
- $f(n)$  must be **lower** than TRUE distance to the goal **via  $n$** 
  - $f(n)$  should be an *honest guess* toward “true” value
  - If  $f(n)$  **over-estimates**, then a minimizer algo may be **wrongly** guided
  - *Admissibility is required in order to guide correctly: for “optimality”*
- Why not use **true** distance for  $f$ ? *Then, who needs a search :)*
  - *Heuristic*  $\equiv$  You are just trying to guess!

## 2) Consistent Heuristic for A\* $\equiv$ Guides only toward correct direction

- For every pair of nodes  $n, n'$ , such that  $n'$  is a child of  $n$  on a path from start to goal node
- **If**, true distance  $d(n, n') \geq h(n) - h(n')$  (*consistent*)
  - Means *Triangle inequality*, or
  - Heuristic function increases *monotonically*:  $h(n') + d(n, n') \geq h(n)$
- **Then**, Guidance is perfect  $\equiv$  Never in a wrong direction  $\equiv$  *Consistent*
  - NO backtrack necessary beyond “*depth contour*” of goal on DFS
    - [*backtracking wastes time*]



*Note: heuristic function is  $f(n) = g(s \text{ to } n) + h(\text{estimated from } n \text{ to goal})$   
and,  $d(n, n') = \text{actual distance in input between } n \text{ and } n'$*

# Consequence of Consistent Heuristic for A\*


- **Consistent  $\equiv$  Guidance is perfect:**
  - will find a path if it exists &**
  - will find a shortest path**
- **Search “flows” like water downhill toward the goal,**
  - following the best path
  - **Flows perpendicular to the contours of equal  $f$ -values**
- *Otherwise,*

// proof: book slides Ch04a sl 30

# Consequence of Consistent Heuristic for A\*

- If, true distance  $d(n, n') \geq h(n) - h(n')$
- Heuristic function increases monotonically:  $h(n) + d(n, n') \geq h(n')$
- Then, Guidance is perfect  $\equiv$  Never in a wrong direction  $\equiv$  Consistent
- NO backtrack necessary on DFS [backtracking wastes time]
- *Otherwise,*

Say,  $h(n) - h(n') = 300 - 270 = 30$ , but  $h(n) - h(n'') = 300 - 240 = 60$   
 $d(n, n'') = 50$ , i.e.  $50 \not\geq 60$ , or heuristic is **not** consistent



## *Optimality of $A^*$*

- The book has a bit of confusion on finding the goal vs. finding “a” best path to the goal
- Admissibility is *necessary*, i.e., without admissibility guidance may be wrongly directed;  
but not *sufficient*: may get into infinite loop (same  $f$ )
- Consistency ( $f$  must increase) implies admissibility as well
- For **consistent** heuristic DFS  
No backtracking is necessary (beyond goal  $f_{goal}$ )
  - Optimality of consistency = if stuck, then no finite path to goal
- HOWEVER, generic search is NP-hard problem, there may be Exponential number of nodes within the contour of the goal



## *Iterative deepening and $A^* = \underline{IDA}^*$*

- Consistent heuristic is **difficult** to have in real life
- **Inconsistent** heuristic needs backtracking on  $A^*$ 
  - Just as we discussed in DFS
- A misguided search (*with inconsistent heuristic*) may get stuck:
  - infinite search
  - So, what to do? Answer: **IDA\***

## *Iterative deepening and $A^* = IDA^*$*

- IDS +  $A^*$  provides *guarantee* of finding solution:  $IDA^*$  at the cost of repeated run of  $A^*$
- Do not let  $A^*$  go to arbitrary depth
- $IDA^*$ : fixed depth  $A^*$ ,
  - increase **depth** (by  $f$ ) incrementally and **restart**  $A^*$
  - Note: backtrack is allowed on  $A^*$   
but no need for that with consistent heuristic

## *Recursive Best-first Search*

- *It is still by,  $f = g + h$*
- Remember **second best  $f$**  of some ancestors
- If current  $f$  ever becomes **>** that (say, node  $s = \text{best-}f\text{-so-far}$ ),
  - then jump back to that node  $s$  (Fig 3.27)
- Needs **higher memory** than DFS (more nodes in memory), but a **compromise** between DFS and BFS

## *Forgetful search: SMA\**

- Often **memory** is the main problem with A\* variations
- *Memory bounded search:*
  - **Forget** some past explored nodes,
    - Some similarity with depth-limited search
- **SMA\*** optimizes memory usage
- A\*, but **drops worst  $f$  node** in order to expand, when memory is full
  - However, the  $f$  value of the forgotten node is not deleted
  - It is kept on the parent of that node
  - When current  $f$  crosses that limit,  
it restarts search from that parent (*Simplified MA\**)
- SMA\* is **complete**: goal is found always,  
**provided** start-to-goal path can fit in memory (*no magic!*)

## *Real life AI search algorithms*

- Two more strategies in search algorithms are prevalent,
  - in addition to these search algorithms
- 1. *Learning*: Even from early AI days
  - Learn how to play better, success/failure as training
    - To **learn**  $h$  function, or other strategies
- 2. *Pattern Database*:
  - Human experts use patterns in memory!
  - IBM DeepBlue series stored and matched successful patterns
  - Google search engine evolved toward this
  - Database **indexing** becomes the key problem, for fast matching

## *Summary of informed search*

- Best first search uses some  $f(\text{node})$  to guide search
- $f(n) = g(n)$ , actual distance on the path from source
  - search is on circular contour
- $f(n) = g(n) + h(n)$  → A\* search
- Admissible heuristic,  $h(n) \leq h^*(n)$ , actual optimum
  - A\* is complete, optimal,  
but may include suboptimal path to goal,  
however, it will not stop until optimal path is found
  - A\* follows narrowing contour, until C\* contour but no more
- Consistent,  $h(n) \leq h(n_1) + d(n, n_1)$ , if  $n$  is expanded before  $n_1$ 
  - Euclidean triangle inequality is preserved
  - No suboptimal path to goal is explored
  - Less node explored, most efficient situation
- Consistent means admissible, but not the other way round

*Code A\* search and run on the Romanian road problem, or your favorite search problem*

*Our Next Module: Local Search*

# MORE SEARCH ALGORITHMS

- Local search: Only achieving a *goal* is needed,
  - not the path to the goal
- Possibilities:
  - Non-determinism in state space (graph)
  - Partially-observable state space
  - On-line search (full search space is not available)



# LOCAL SEARCH ALGORITHMS

- Only goal is needed, NOT the path (*8-puzzle*)
- Sometimes: No goal is provided, only how to do “better”
- Algorithm is allowed to forget past nodes
- Search on **objective function** space ( $f(\underline{x})$ ,  $\underline{x}$ )
  - $\underline{x}$  , node / state / location / coordinate, may be a vector
  - $f$  is continuous or discrete,
    - $f$  continuous means direction of betterment is available as *Grad(f)*
- Examples:
  - $n$ -queens problem (on chess board, no queen to attack another)
    - Search space: a node –  $n$  queens placed on  $n$  columns
  - VLSI circuit layout
    - (possibly!)  $n$  elements connected, optimize total area
      - Or, given a fixed area, make needed connections

# LOCAL SEARCH ALGORITHMS

- Hill-climbing search → *A Greedy Algorithm*
- Take the best move out of all possible **current** moves
- 8-queens problem: minimize  $f = \# \text{ attacks}$ : *Fig 4.3*
  - Move a queen on its column, 8x7 next moves (8q, 7col to try for each)

# LOCAL SEARCH ALGORITHMS

- Representation may be important for efficiency
  - Queen positioned at  $(x, y)$ ?  
~  $8 \times 8 = 64$  possibilities:  $64 \times 8$  search space
  - Alternative: Queen is linked to a column (cannot have more than one queen per column) and its position is row#  $y$ ,  
7 possibilities only:  $7 \times 8$  search space
- Complexity may depend on representation

# LOCAL SEARCH ALGORITHMS

- Consequence of forgetting past: may get stuck at a local minimum
  - If problem is *NP-hard*, typically exponential # local minima
- Local peak, ridge, local plateau, shoulder (max-problem):
  - No better move available, but..
  - There may exist better maxima after one/more valley
- Solution? Jump out of local minima, **random move**
  - 8-queens problem: **random restart the algorithm**
- How do you know your next optimum is better than the last one?
  - Why do you have to forget everything?
  - Just remember the last best optimum and update if necessary!
  - Still not guaranteed to find <<Global>> optimum, but  $f^*$  converges

## LOCAL SEARCH ALGORITHMS: Simulated Annealing

- *Systematic* Random move
- If improvement  $\text{delta-}e \equiv d > \text{threshold}$ , then make the move
  - Otherwise, ( $d \leq \text{threshold}$ ), generate a random number  $r$  and
  - If  $r \geq \exp(d/T)$ , then take the “wrong” move anyway
    - Otherwise, **randomly** jump to a node
- $T$ : a constant, but from a sequence of reverse sort
  - Reduces in each iteration: *less and less random restarts*
    - E.g., 128, 64, 32, 16, 12, 8, 6, 4, 3, 2, 1
- Concept comes from **metallurgy**/ condensed matter physics:
  - slowly reduce temp  $T$  from high to low:
    - molecules settle to **lowest energy = strongest bonding**

## LOCAL SEARCH ALGORITHMS: Beam search

- Keep  $k$  best nodes in memory, and do Hill-climbing from each
  - Same time!
- Independent  $k$  nodes, but they are **best  $k$  nodes so far**
- Still has a tendency to converge to a local minimum, unless
  - $k$  initial samples are “good” samples in statistical sense
- Some form of **evolution** is taking place!
  - With  $k$  children producing  $k$  offspring
    - Why not formalize it? → *Stochastic Beam Search*
    - *Stochastic Beam Search: Choose  $k$  offspring probabilistically*

# LOCAL SEARCH ALGORITHMS: Genetic Algorithm: Fig. 4.6-7

- Further formalization of Stochastic Beam Search from Biology
  - Genetic Algorithm
- Keep independent  $k$  (even number) nodes in memory
- Pair up, and crossover  $\rightarrow$  two new nodes
  - Needs special representation of states/nodes: see 8-q
  - Crossover position is random
- Participation in reproduction is based on objective function  $f$ :
  - Not all  $k$  nodes participate in the crossover, some are culled
  - Higher  $f$  means higher probability of participation
- Random mutation between iterations are allowed
  - With a very low probability (infrequently)
  - To jump out of any local optima
- *Many types of variation of GA's exist, on each of the above choices*

# CONTINUOUS SPACE: LOCAL SEARCH

- Numerical optimization in Math
  - *a 200 year old subject*
- $(f(\underline{x}), \underline{x})$ :  $\underline{x}$  is continuous, and  $f$  is continuous differentiable
- Infinite search space (not a graph search), even if bound
  - (infinite states if  $\underline{x}$  is real)
- SA is useful in continuous space, but not GA
  - GA needs discrete representation
  - Discretization of  $\underline{x}$  is some-times used in GA for continuous space
- Local search:
  - Typically, start coordinate and the resulting path is not important
- Main help comes from the existence of **derivative of  $f$** :
  - Second or higher order derivatives may be needed:
    - *analytic  $f$ : Cauchy-Riemann condition helps*



# CONTINUOUS SPACE: LOCAL SEARCH

## Optimization theory in Math

- Gradient provides **direction** of next move,
  - $\underline{x} \leftarrow \underline{x} + a * (-\text{sign\_of}(\text{Del}[f(\underline{x})])), f$  is the objective function
  - but, how far? What value of  $a$ ?
  - Many algorithms exist for finding  $a$
  - E.g., “line search” algorithm
- Newton-Raphson:
  - 1D:  $x \leftarrow x - f(x)/f'(x)$ , if analytic solution exists
  - nD: derived by using  $\text{gradient}(f) = 0$ 
    - $a = \text{Hessian}, H^{-1}(f)$ ,
      - // a partial-double derivative matrix,  $H_{ij} = d^2 / dx_i dy_j$
- Conjugate gradient: faster convergence
  - better optimization for longer vision ahead
  - winds down along the valleys

- Typically, optimization needs **constraints**:
  - Limit the search space!
- *Linear programming*, with inequality constraints:
  - Search space is a closed polygon (lines enclosed by lines)
  - *Simplex* algorithm, *Interior-point* search algorithm
- *Quadratic Programming*
- General case: Machine learning
  - Support vector machine
  - Neural network
  - ...

# NON-DETERMINISTIC SEARCH SPACE

- Typically, discrete, or else ☹
- If-else on nodes,
  - condition dependent
  - You do not know what will show up on a node – plan ahead!
  - *Example: Mars rover*
- Not as bad as it seems: *And-Or tree*
- Deterministic search: ‘Or’ graph, chose a node or the other
- Non-deterministic search: All nodes may need to be looked into
  - **And-Or tree** search
  - ‘And’ on if-else node: take all options in search: *Fig 4.10*
- Non-determinism may come from failure from a move
  - Loops are possible → *try and try again, until give up!*

# PARTIALLY OBSERVABLE SEARCH SPACE

- *Unobservable* = completely blind = unknown search space, really! ☹
  - Not as bad as it sounds: just get the job done! *Fig. 4.9*
- *Partially observable*: think of a blind person with the stick!
  - Sensor driven search
- State space: beliefs on which nodes (a set) the agent may be in
  - *A node does not know its own exact id*
  - *Discrete space, please!*
- Action reduces possibilities or nodes, ..
  - Action may be exploratory, just to reduce uncertainty of possible nodes the agent is in
  - Converging (fast) to the goal node → you are sure you attained it!
  - See the vacuum cleaner example:
    - A few moves blindly → room is cleaned

# “ON-LINE” SEARCH

- Off-line: Search space is known, at least static
  - Non-deterministic: at least possibilities are known or “static”
- On-line: Robots
  - Compute-act-sense-compute-.....
  - State space develops as-we-go:
    - not “thinking” algorithm, actually make the move, no way to un-commit
  - Not much different
- Strategy, heuristics,
  - but no set of nodes, may be finite/infinite